

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from imblearn.over_sampling import SMOTE
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
from imblearn.over_sampling import SMOTE

```

Data Loading

```

df = pd.read_csv(r"C:\Users\arjun\Downloads\archive (3)\spam.csv",
encoding="ISO-8859-1") # Try this first
df.head()

```

	v1	v2	Unnamed: 2
0	ham	Go until jurong point, crazy.. Available only ...	NaN
1	ham	Ok lar... Joking wif u oni...	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN
3	ham	U dun say so early hor... U c already then say...	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN

	Unnamed: 3	Unnamed: 4
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN

Data Preprocessing

```

df.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], axis=1, inplace=True)

```

```
df.head()

   v1 v2
0 ham Go until jurong point, crazy.. Available only ...
1 ham      Ok lar... Joking wif u oni...
2 spam Free entry in 2 a wkly comp to win FA Cup fina...
3 ham U dun say so early hor... U c already then say...
4 ham Nah I don't think he goes to usf, he lives aro...

df['v1'].value_counts()

v1
ham      4825
spam      747
Name: count, dtype: int64

df['v1']=df['v1'].map({'ham': 0, 'spam': 1})

df.head()

   v1 v2
0  0 Go until jurong point, crazy.. Available only ...
1  0      Ok lar... Joking wif u oni...
2  1 Free entry in 2 a wkly comp to win FA Cup fina...
3  0 U dun say so early hor... U c already then say...
4  0 Nah I don't think he goes to usf, he lives aro...
```

Model Training

```
vectorizer = TfidfVectorizer(stop_words='english')
X = vectorizer.fit_transform(df['v2']) # Replace 'message' with your
text column
y = df['v1']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42, stratify=y)

# Apply SMOTE to balance classes
smote = SMOTE(random_state=42)
X_train_bal, y_train_bal = smote.fit_resample(X_train, y_train)

# Train Naïve Bayes classifier
nb = MultinomialNB()
nb.fit(X_train_bal, y_train_bal)

# Make predictions
y_pred = nb.predict(X_test)
```

```

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
print("Classification Report:\n", classification_report(y_test,
y_pred))


# Confusion matrix visualization
plt.figure(figsize=(5, 4))
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d',
cmap='Blues', xticklabels=['Ham', 'Spam'], yticklabels=['Ham',
'Spam'])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()

```

Accuracy: 0.97

Classification Report:

	precision	recall	f1-score	support
0	0.99	0.98	0.98	966
1	0.87	0.93	0.90	149
accuracy			0.97	1115
macro avg	0.93	0.96	0.94	1115
weighted avg	0.97	0.97	0.97	1115



```

# Training Accuracy
train_accuracy = nb.score(X_train_bal, y_train_bal)
print(f"Training Accuracy: {train_accuracy:.2f}")

# Testing Accuracy
test_accuracy = nb.score(X_test, y_test)
print(f"Testing Accuracy: {test_accuracy:.2f}")

# Compare Train & Test Accuracy
if train_accuracy - test_accuracy > 0.1: # 10% difference is a sign
of overfitting
    print("Potential Overfitting Detected!")
else:
    print("No Overfitting Detected.")

Training Accuracy: 0.99
Testing Accuracy: 0.97
No Overfitting Detected.

```

Testing Model

```

sample_text = ["Congratulations! ☑ You have won a FREE iPhone 15!
Click the link below to claim your prize now. Hurry, limited time
offer! ☑ http://fake-prize.com"]

# Convert to numerical features (Assuming you used TF-IDF or
CountVectorizer)
sample_text_vectorized = vectorizer.transform(sample_text)

# Predict
prediction = nb.predict(sample_text_vectorized)

# Output Result
print("Spam" if prediction[0] == 1 else "Not Spam")

Spam

# Example message to test (Non-Spam)
test_message = ["Hey, are we still meeting for lunch at 1 PM? Let me
know if you're running late!"]

# Convert text into numerical features (assuming you used TF-IDF or
CountVectorizer)
test_message_transformed = vectorizer.transform(test_message) # Use
the same vectorizer from training

# Predict using the trained model
prediction = nb.predict(test_message_transformed)

```

```
# Output result  
print("Spam" if prediction[0] == 1 else "Not Spam")
```

Not Spam

Detailed Report